

MPCS 51040 – C Programming

Lecture 10 – Optimization & Parallel Programming

Dries Kimpe

November 28, 2016



General Info

- ▶ HW6 grading in progress
- ▶ Quiz: see repository
- ▶ Project



Project



Project

Any questions/help needed?

Reminder:

- ▶ Review early
- ▶ Clarifications/questions: use piazza!



Profiling

What is profiling?

*In software engineering, profiling (“program profiling”, “software profiling”) is a form of dynamic program analysis that measures, for example, the space (memory) or time complexity of a program, the usage of particular instructions, or the frequency and duration of function calls. Most commonly, **profiling information serves to aid program optimization.***

Profiling is achieved by instrumenting either the program source code or its binary executable form using a tool called a profiler (or code profiler).

Profilers may use a number of different techniques, such as event-based, statistical, instrumented, and simulation methods.

What is the relationship to parallel programming?

- ▶ We want to direct our efforts so to have the most impact; In other words, we want to speed up the part of the code where most of the processing time is spent.
- ▶ Profiling is (one of) the methods to find those code areas.



Profiling

gprof

How does it work?

- ▶ Compiler instruments your code to collect information **at runtime** about which functions are being called and by whom.
Note: causes (limited) slowdown and disturbs program.
- ▶ At certain intervals, **during the execution of your program** the current *program counter* is examined to determine which *function* is currently being executed.

Some caveats:



- ▶ We use gprof because it is readily available – it is not necessarily the best tool for the job. . .
- ▶ gprof is not well suited for profiling multi-threaded programs.
- ▶ Compiler optimization can interfere (see later).



<https://sourceware.org/binutils/docs/gprof/>



```

today% gprof game gmon.out
Flat profile:

Each sample counts as 0.01 seconds.
%   cumulative   self           self      total
time  seconds    seconds   calls   s/call   s/call   name
58.63    6.26    6.26 97951385    0.00    0.00  searchDir
19.84    8.37    2.12 331558518    0.00    0.00  get
9.00    9.40    1.03  8894885    0.00    0.00  board_has_winner
4.88    9.92    0.52    9    0.06    1.17  rate_board
1.89   10.10    0.18  8894866    0.00    0.00  board_unplay
1.31   10.25    0.14  8894875    0.00    0.00  board_play
1.13   10.37    0.12 17789741    0.00    0.00  set
0.94   10.47    0.10 17620390    0.00    0.00  board_can_play_move
0.89   10.56    0.10    1    0.00    0.00  main
0.61   10.63    0.07    10    0.01    0.01  board_destroy
0.33   10.66    0.04    18    0.00    0.00  board_can_play
0.09   10.67    0.01  8894886    0.00    0.00  board_get_width
0.09   10.68    0.01  4405075    0.00    0.00  other_player
0.00   10.68    0.00    13    0.00    0.00  player2string
0.00   10.68    0.00    11    0.00    0.00  agent_describe
0.00   10.68    0.00    11    0.00    0.00  board_get_height
0.00   10.68    0.00    11    0.00    0.00  board_print
0.00   10.68    0.00    11    0.00    0.00  computer_agent_describe
0.00   10.68    0.00    9    0.00    1.17  agent_play
0.00   10.68    0.00    9    0.00    0.00  board_duplicate
0.00   10.68    0.00    9    0.00    1.17  computer_agent_play
0.00   10.68    0.00    9    0.00    0.00  next_player
0.00   10.68    0.00    3    0.00    0.00  readnum
0.00   10.68    0.00    2    0.00    0.00  agent_destroy
0.00   10.68    0.00    2    0.00    0.00  computer_agent_create
0.00   10.68    0.00    2    0.00    0.00  computer_agent_destroy
0.00   10.68    0.00    2    0.00    0.00  new_computer_player
0.00   10.68    0.00    2    0.00    0.00  select_player
0.00   10.68    0.00    1    0.00    0.00  board_create
0.00   10.68    0.00    1    0.00    0.00  create_board

```

How to use it?

1. Compile and link with `-pg` (in addition to normal options)
2. Run the program as usual. Program must exit 'cleanly'.
3. At exit, file `gmon.out` is written in the **current directory**.
4. Analyze `gmon.out` using the `gprof` tool.
`gprof <programexecutable> <gmonfile>`



- ▶ Compile/ link hw6
- ▶ Generate profile
- ▶ View using `gprof`
- ▶ Demonstrate on optimized binary

Analyzing GProf Output

Flat Profile

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
58.63	6.26	6.26	97951385	0.00	0.00	searchDir
19.84	8.37	2.12	331558518	0.00	0.00	get
9.66	9.40	1.03	8894885	0.00	0.00	board_has_winner
4.88	9.92	0.52	9	0.06	1.17	rate_board

- ▶ Shows how much time was spent in each function (total), in descending order.
- ▶ Order by time *not inclusive* of time spent in children (i.e. other functions called by this function)
 - see *cumulative* for inclusive time.



<https://sourceware.org/binutils/docs/gprof/Flat-Profile.html>



Analyzing GProf Output

Call Graph

Shows each function, which functions it called and which functions called it.

			8894866	rate_board [4]	
		0.52	9.96	9/9	computer_agent_play [3]
[4]	98.2	0.52	9.96	9+8894866	rate_board [4]
		1.03	8.32	8894875/8894885	board_has_winner [5]
		0.18	0.12	8894866/8894866	board_unplay [8]
		0.14	0.06	8894866/8894875	board_play [9]
		0.10	0.00	17620300/17620300	board_can_play_move [11]
		0.01	0.00	4405075/4405075	other_player [15]
		0.01	0.00	8894875/8894886	board_get_width [14]
			8894866	rate_board [4]	



<https://sourceware.org/binutils/docs/gprof/Call-Graph.html>



C11 and Parallel Programming

C11 describes a header (`threads.h`) which provides functions for creating, destroying and synchronizing threads.

Thread support is *optional* in C11, and most compilers do not currently provide C11 thread functionality.

- ▶ This is mostly since there have been non-C11 standard specified methods of using threads. For example: POSIX threads (pthreads) described by POSIX.1c standard are available on most common platforms (including Windows, Linux, OS X, Solaris, OpenBSD)

We will be using pthreads (`pthreads.h`)



Review 'man pthreads', in particular the section 'Thread-safe functions'.



POSIX Threads

Example program

```
1 void * PrintHello(void *threadid)
2 {
3     long tid;
4     tid = (long)threadid;
5     printf("Hello World! It's me,"
6           "thread %ld!\n", tid);
7     pthread_exit(NULL);
8 }
9 int main(int argc, char *argv[])
10 {
11     pthread_t threads[NUM_THREADS];
12     int rc;
13     long t;
14     for(t=0;t<NUM_THREADS;t++)
15     {
16         rc = pthread_create(&threads[t],
17                             NULL, PrintHello, (void *)t);
18     }
19     // need to join before exiting...
20 }
```

- ▶ Explicitly create and destroy threads.
- ▶ Synchronization through mutex, barrier, condition variables.
- ▶ Communication via Shared memory – threads can access each others variables.
- ▶ Low level interface.
- ▶ Supports both MPMD and SPMD.



Thread Creation/Destruction

```
1  ...
2  for (int i=0; i<count; ++i)
3  {
4      pthread_create(&threads[i],
5                    NULL, printHello, (void *) (intptr_t) i);
6  }
7
8  // Join all threads
9  for (int i=0; i<count; ++i)
10 {
11     // — BLOCKING — operation
12     void * returncode;
13     pthread_join(threads[i], &returncode);
14 }
15 ...
```

- ▶ Program starts out with a single thread
- ▶ We manually (explicitly) create additional threads
 - ▶ `pthread_create` does not **block** or need to wait for the new thread to start executing – but might.
 - ▶ Creating a thread is not a cheap operation.
- ▶ Threads end when calling `pthread_exit` or returning from 'initial' function.
- ▶ *joining* a thread allows waiting for a thread to exit (and obtaining a return code from the thread). The call **blocks** until the indicated thread has finished.



See `pt1.c` example program

- ▶ Compile and link (`-pthread`)
- ▶ Run with multiple thread counts and observe scheduling



Threads

Thread-safe functions (2)

```
1 void * found;  
2 bool find_node(node_t * node, void * d)  
3 {  
4     if (!node)  
5         return false;  
6     if (node->data==d)  
7     {  
8         found=node;  
9         return true;  
10    }  
11    else if (find_node(node->left, d))  
12    {  
13        return true;  
14    }  
15    else  
16        return find_node(node->right, d));  
17 }  
  
1 int pseudorandom()  
2 {  
3     static_assert(sizeof(int)>=4,  
4         "proper size int");  
5     #define M 2147483647
```

For a function to be thread-safe, at the very least, any access to a *shared* resource (memory or other) should be properly synchronized/protected.

Examples:

- ▶ Concurrent access to **the same** variable.
- ▶ Concurrent output to a file or to the screen.

Are these examples thread-safe? Why/Why not?



Thread-safety

Problem Example

```
1  ...
2  unsigned int total = 0;
3
4  void * printHello (void * threadid)
5  {
6      ++total;
7      // or pthread_exit(NULL);
8      return NULL;
9  }
10 ...
```

Question

- ▶ What is the problem here?
- ▶ How to fix it? (see next slide)



See pt2.c



Mutex

Limiting concurrent access

```
1  ...
2  unsigned int total = 0;
3
4  pthread_mutex_t mutex =
5      PTHREAD_MUTEX_INITIALIZER;
6
7  void * printHello (void * threadid)
8  {
9      pthread_mutex_lock(&mutex);
10     ++total;
11     pthread_mutex_unlock(&mutex);
12     return NULL;
13 }
14 ...
```



Mutexes **limit**
concurrency!

(Remember Amdahl's law?)

A *mutex* allows us to have *mutual exclusion*

- ▶ Only one thread can **hold** (or lock) a mutex at a time.
- ▶ Other threads trying to lock the mutex will **block** (wait) until the mutex is unlocked.
- ▶ Only a single waiting thread will succeed in locking the mutex

There are many different kinds of mutexes; They differ on:

- ▶ Fairness
- ▶ Recursive locking
- ▶ Type (reader/writer)
- ▶ Use case (expected contention)



See pt3.c

